

Statistical Methods for Machine Learning

Assignment 2: Basic Learning Algorithms

Philip Pickering
pgpick@gmx.at

Marco Eilers
eilers.marco@googlemail.com

Thomas Bracht Laumann Jespersen
ntl316@alumni.ku.dk

1 Regression

1.1 Maximum Likelihood solution

Use linear model

$$y(\mathbf{x}, \mathbf{w}) = w_0 + w_1x_1 + w_2x_2 + \cdots + w_Dx_D$$

and for the D variables we let $\phi_i(\mathbf{x}) = x_i$ for $i = 1, \dots, D$ and $\phi_0(\mathbf{x}) = 1$.

1.1.1 Selection 1

For our first selection S_1 our design matrix becomes a 200×5 matrix.

$$\Phi_{S_1} = \begin{bmatrix} 1 & \mathbf{x}_{1,1} & \mathbf{x}_{1,2} & \mathbf{x}_{1,3} & \mathbf{x}_{1,4} \\ & & \vdots & & \\ 1 & \mathbf{x}_{i,1} & \mathbf{x}_{i,2} & \mathbf{x}_{i,3} & \mathbf{x}_{i,4} \\ & & \vdots & & \\ 1 & \mathbf{x}_{N,1} & \mathbf{x}_{N,2} & \mathbf{x}_{N,3} & \mathbf{x}_{N,4} \end{bmatrix}$$

where the notation $\mathbf{x}_{i,j}$ indicates the j 'th entry in the i 'th vector.

Finding the ML estimate of our parameters for S_1 gives

$$\mathbf{w}_{S_1} = \begin{bmatrix} -43.0947 \\ -0.1299 \\ 0.0352 \\ 0.9335 \\ -0.0433 \end{bmatrix} \quad \text{and} \quad \text{RMS}_{S_1} = 4.3897$$

1.1.2 Selection 2

Our second selection S_2 consists only of the data from the 'Abdomen 2' column, giving a design matrix Φ_{S_2} of dimensions 200×2 . Training the model on the same training data yields:

$$\mathbf{w}_{S_2} = \begin{bmatrix} -37.4085 \\ 0.6133 \end{bmatrix} \quad \text{and} \quad \text{RMS}_{S_2} = 5.2064$$




Figure 1: Plot of RMS against varying values of α

1.1.3 Discussion

Just looking at the root mean square values of the two selections, it appears that S_1 performs better than S_2 , but not by a lot. This could suggest that either the variable ‘Abdomen 2’ is the most descriptive in terms of body fat, or that the linear model simply is a poor fit no matter how many variables we include. It could be a combination of the two.

It could probably be argued that the linear regression model is a poor predictor, but including more variables should improve the results.

1.2 Maximum a posteriori solution

For S_1 the lowest root mean square value is 4.3614, the same as the ML solution differing on the second digit. The difference is larger for S_2 in which we obtain the value 5.1071, differing from the ML solution on the first digit.

In fig. 1 two plots are found of the root mean square values for varying values of α . In both plots we set $\beta = 1$. The RMS value of the ML solution is plotted as a straight line.

1.2.1 Comparison

Firstly, we can observe for both plots that when $\alpha = 0$, we obtain the same RMS error for the MAP estimate as for the ML solution. This is expected and demonstrates that when our prior precision parameters are set to zero, the MAP estimate becomes the ML estimate.

Fig. 1(a) is the plot for S_1 , and it can be seen that the RMS_{MAP} error drops below the RMS_{ML} in the interval $[-0.091, 0]$. In fig. 1(b) the plot for S_2 similarly gives us that the RMS_{MAP} error is lower in the interval $[-0.844, 0]$.

The intervals in which the MAP estimate outperforms the ML estimate are very small, and the obtained difference likewise small.

1.3 Theory

Verify result in equation (3.49) for the posterior distribution of the parameters \mathbf{w} in the linear basis function in which \mathbf{m}_N and \mathbf{S}_N are defined




Figure 2: Visualisation of the training data for each of the KNOLL problems

	Training	Test
KnollA	1%	3%
KnollB	40%	49%
KnollC	1%	3%

Table 1: Error percentages for each of the KNOLL problems

2 Linear Discriminant Analysis

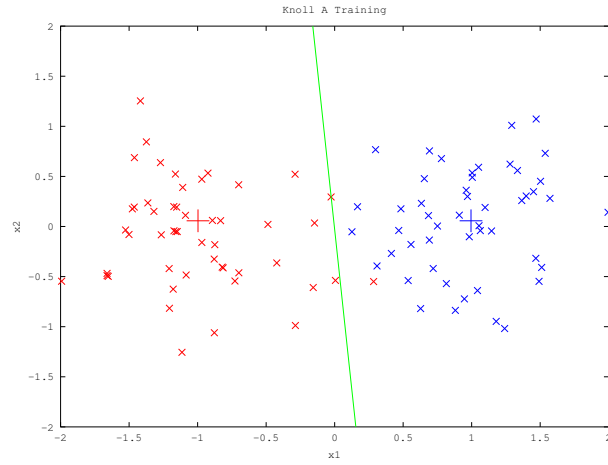
2.1 Observations regarding the data

In fig. 2 you can see plots of all three training data sets. It is obvious that while KnollA and KnollC contain well-separable groups of data points, the distributions in KnollB seem to overlap. We can therefore expect that LDA will perform relatively well on KnollA and KnollC, whereas the results for KnollB will likely be more erratic.


2.2 Results of LDA

If we run LDA on all three data sets and plot them with their respective decision boundaries (fig. 3), we find that the algorithm indeed performs well on the A and C sets. In table 1 we can see that although the error count is greater than zero for all training and test sets, which is also implied by the margin always being slightly negative, the number of errors is relatively low for both training and test sets. The training set performs slightly better than the test set in both cases, which is to be expected, but there is not over- or underfitting problem.


For KnollB, however, the results are much worse. The properties of the underlying distribution simply do not allow them to be separated by a straight line, which results in a high error percentage and a highly negative margin. Since the distributions in B overlap to a high degree, it is generally not possible to find a function of any kind to separate both classes accurately. However, a Quadratic Discriminant Analysis might be able to find out that one class concentrates a lot more in a certain range, whereas the other one has a higher variance and is therefore spread out more widely. It could therefore assign all points in this range to one class and the rest to the other. While this will still result in a high error percentage, such an algorithm might actually find the Bayes Optimal Solution.




(a) as




(b) df




(c) as



(d) df



(e) as



(f) df

Figure 3: Visualisation of the training data for each of the KNOLL problems

k	KNOLLA		KNOLLB		KNOLLC	
	test	train	test	train	test	train
1	0.04	-	0.23	-	0.20	-
3	0.02	0.01	0.20	0.13	0.32	0.16
5	0.03	0.01	0.19	0.19	0.34	0.21
7	0.04	0.01	0.19	0.19	0.42	0.21
9	0.03	0.02	0.23	0.20	0.46	0.33
	3	3	5	3	1	3

Table 2: Table of results for k -NN classifier on all three KNOLL problems.

3 Nearest Neighbor Classification

3.1 Nearest Neighbor Classification with Euclidian Metric

The results of the runs of the classifier for different values of k are found in table 2. The bottom field of each column indicates the k for which the classifier gave the lowest error rate.

The error rate is computed by adding up the errors made by the classifier, and then dividing by the test set size. As we have the actual class of a given point (for both test and training sets), this is straightforward to implement.

As a general note, classifying the training set itself will always give us $k = 1$ as the optimal value, which makes sense, because the closest point to a point in both the training and test sets is the point itself, hence the distance will always be zero (and the classification perfect). But it's interesting to see how the classifier behaves when more than one neighbor has to be considered. If a point belonging to one class is surrounded by points in the other class, then it will be misclassified. As can be seen from the results, the optimal value for k in all the training sets is 3, although it should be noted that the error rate for $k = 3, 5$ and 7 in KNOLLA are all 0.1.

In KNOLLA, the error rates are very low (< 0.05) for all the values of k , being lowest when $k = 3$. In KNOLLB and KNOLLC the error rates increase dramatically (≥ 0.19), for KNOLLB the lowest error rates are obtained when $k = 5$ or 7. In KNOLLC the best results are unanimously obtained when $k = 1$.

Observations

We can compute an accuracy for the k -NN on the same training data, which for $k = 1$ always gives an error degree of zero, because the nearest neighbour of a given point in the data set is the point itself. But for higher values of k , we can get errors for points surrounded by points from the opposite class.

3.2 Changing the Metric

To prove that d is a metric, given

$$d(\mathbf{x}, \mathbf{z}) = \|\mathbf{M}\mathbf{x} - \mathbf{M}\mathbf{z}\|, \text{ where } \mathbf{M} = \begin{pmatrix} 100 & 0 \\ 0 & 1 \end{pmatrix}$$

and $\|\cdot\|$ is the standard L_2 -norm (in \mathbb{R}^2), we need to verify $\forall \mathbf{x}, \mathbf{y} \in \mathbb{R}^2$ that 1) $d(\mathbf{x}, \mathbf{y}) \geq 0$; 2) $d(\mathbf{x}, \mathbf{y}) = 0 \Leftrightarrow \mathbf{x} = \mathbf{y}$; 3) $d(\mathbf{x}, \mathbf{y}) = d(\mathbf{y}, \mathbf{x})$ (symmetry) and 4) $\forall \mathbf{x}, \mathbf{y}, \mathbf{z} \in \mathbb{R}^2 : d(\mathbf{x}, \mathbf{z}) \leq d(\mathbf{x}, \mathbf{y}) + d(\mathbf{y}, \mathbf{z})$.

3.2.1 Proof

We need only to observe that \mathbf{M} is a projection $m : \mathbb{R}^2 \rightarrow \mathbb{R}^2$, i.e. onto \mathbb{R}^2 itself, given by $m(\mathbf{x}) = \mathbf{M}\mathbf{x}$. This immediately gives us all the properties we need, because L_2 is itself a (complete) metric on \mathbb{R}^2 .

For instance, if we let $\mathbf{x}, \mathbf{y}, \mathbf{z} \in \mathbb{R}^2$ and $\mathbf{x}', \mathbf{y}', \mathbf{z}'$ be the result of applying m on $\mathbf{x}, \mathbf{y}, \mathbf{z}$ respectively, we can prove the triangle inequality:

k	training set	test set
1	0.04	-
3	0.02	0.02
5	0.03	0.03
7	0.04	0.04
9	0.03	0.03
	3	3

$$\begin{aligned}
d(\mathbf{x}, \mathbf{z}) &= \|\mathbf{M}\mathbf{x} - \mathbf{M}\mathbf{z}\| \\
&= \|\mathbf{M}\mathbf{x} - \mathbf{M}\mathbf{y} + \mathbf{M}\mathbf{y} - \mathbf{M}\mathbf{z}\| \\
&= \|(\mathbf{x}' - \mathbf{y}') + (\mathbf{y}' - \mathbf{z}')\| \\
&\leq \|\mathbf{x}' - \mathbf{y}'\| + \|\mathbf{y}' - \mathbf{z}'\| \quad (\text{by property of } L_2 \text{ norm}) \\
&= d(\mathbf{x}, \mathbf{y}) + d(\mathbf{y}, \mathbf{z})
\end{aligned} \tag{1}$$

3.2.2 Results

Using d as metric on KNOLLC we obtain the following results.

The error rates drop dramatically. We could visualize with a plot?